



Deep Learning

Dr. Mehran Safayani

safayani@iut.ac.ir

safayani.iut.ac.ir



<https://www.aparat.com/mehran.safayani>



https://github.com/safayani/deep_learning_course



Department of Electrical and computer engineering, Isfahan university of technology, Isfahan, Iran

Supervised Learning

Input(x)	Output (y)	Application
Home features	Price	Real Estate
Ad, user info	Click on ad? (0/1)	Online Advertising
Image	Object (1,...,1000)	Photo tagging
Audio	Text transcript	Speech recognition
English	Chinese	Machine translation
Image, Radar info	Position of other cars	Autonomous driving

Basics of Neural Network Programming

Binary Classification

Binary Classification



→ 1 (cat) vs 0 (non cat)

			Blue			
	Green		255	134	93	22
Red			255	134	202	22
						2
		255	231	42	22	4
						30
		123	94	83	2	192
						124
		34	44	187	92	34
						142
		34	76	232	124	94
		67	83	194	202	

Binary classification

		Blue				
	Green		255	134	93	22
Red		255	134	202	22	2
		255	231	42	22	4
		123	94	83	2	192
		34	44	187	92	34
		34	76	232	124	94
		67	83	194	202	

$$\vec{x} = \begin{bmatrix} 255 \\ 231 \\ \vdots \\ 254 \\ 253 \\ 250 \\ 220 \end{bmatrix} \quad 64 \times 64 \times 3 = \underbrace{12288}_{n_x}$$

$\vec{x} \rightarrow \text{model} \rightarrow \hat{y}$

- Notation

$$(\vec{x}, y) \quad x \in R^{n_x}, y \in \{0,1\}$$

Binary classification

- m training example: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

• $X = \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & & | \end{bmatrix}$

$\xleftarrow{\quad m \quad} \quad \quad \quad \xrightarrow{\quad n_x \quad}$

$$X \in \mathbb{R}^{n_x \times m}$$

$$X.\text{shape} = (n_x, m)$$

$$Y = [y^{(1)}, y^{(2)}, \dots, y^{(m)}]$$

$$Y \in \mathbb{R}^{1 \times m}$$

$$Y.\text{shape} = (1, m)$$

Logistic Regression

- Given x , output $\hat{y} = P(y=1|x)$ $0 \leq \hat{y} \leq 1$


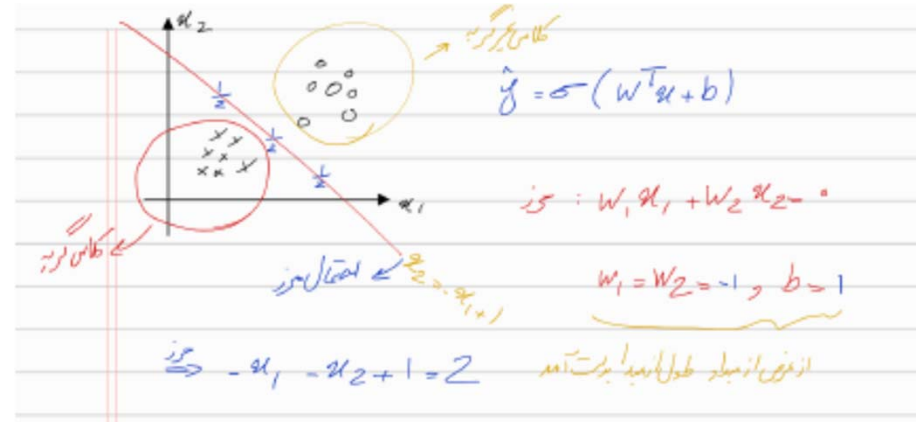
$x \in \mathbb{R}^{n_x}$ parameters: $w \in \mathbb{R}^{n_x}$, $b \in \mathbb{R}$

$$\hat{y} = w^T x + b$$

Sigmoid function $\sigma(z) = \frac{1}{1+e^{-z}}$

if z large $\sigma(z) \approx 1$

if z large negative $\sigma(z) \approx 0$



Logistic Regression

$$\bullet \hat{y} = \sigma(\underbrace{w^T x + b}_z) \quad \hat{y} = w^T x$$

$$x_0 = 1, x \in \mathbb{R}^{n_x+1}$$

$$W = \begin{bmatrix} b = w_0 \\ w_1 \\ w_2 \\ w_3 \\ \vdots \\ \vdots \\ \vdots \\ w_{n_x} \end{bmatrix} \begin{matrix} \left. \vphantom{\begin{bmatrix} b = w_0 \\ w_1 \\ w_2 \\ w_3 \\ \vdots \\ \vdots \\ \vdots \\ w_{n_x} \end{bmatrix}} \right\} b \\ \left. \vphantom{\begin{bmatrix} w_1 \\ w_2 \\ w_3 \\ \vdots \\ \vdots \\ \vdots \\ w_{n_x} \end{bmatrix}} \right\} w \end{matrix}$$

Logistic Regression cost function



- Loss (error) function: $L(\hat{y}, y) = \frac{1}{2} (\hat{y} - y)^2 = \frac{1}{2} (\sigma(w^T x + b) - y)^2$ SE: Square Error
- What is the problem?

• .

Convexity

Function $h(u)$ with $u \in X$ is **convex** if for any $u, v \in X$ and for any $0 \leq \lambda \leq 1$ we have:

$$h(\lambda u + (1 - \lambda)v) \leq \lambda h(u) + (1 - \lambda) h(v)$$



<https://mlstory.org/optimization.html>


$$f(\lambda u + (1 - \lambda)v) \leq \lambda f(u) + (1 - \lambda) f(v)$$

Cross Entropy

- $L(\hat{y}, y) = -(y \log \hat{y} + (1-y) \log(1-\hat{y}))$

if $y=1$: $L(\hat{y}, y) = -\log \hat{y}$

if $y=0$: $L(\hat{y}, y) = -\log(1-\hat{y})$



<https://datamonje.com/classification-loss-functions/>

- Cost function:
$$J(w,b) = \frac{1}{m} \sum_{i=1}^m L(y^{(i)}, \hat{y}^{(i)})$$
$$= -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1-y^{(i)}) \log(1-\hat{y}^{(i)})]$$

Gradient Descent

Cost Function

Minimize $J(\mathbf{b}, w_1)$

\mathbf{b}, w_1

If $J(w_1) = (w_1 - 2)^2$

$$\frac{dJ(w_1)}{dw_1} = 0$$



$$\frac{dJ(w_1)}{dw_1} = 2(w_1 - 2) = 0$$



$$w_1 = 2$$

Gradient Descent

Minimize $J(\mathbf{b}, w_1)$

\mathbf{b}, w_1

Minimize $J(\mathbf{b}, w_1, \dots, w_n)$

$\mathbf{b}, w_1, \dots, w_n$

Repeat until convergence: {

For $j=0, \dots, n$

$\mathbf{b} = w_0$

$$w_j = w_j - \alpha \frac{dJ(\mathbf{b}, w_1, \dots, w_n)}{dw_j}$$

}

α is **learning rate**

Updating all w_j *Simultaneously*

Convergence condition:

$$\|W^{t+1} - W^t\|_2 \leq \varepsilon$$

Gradient Descent

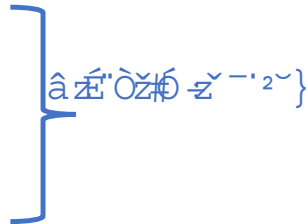
Correct form

$$\text{temp0} = \mathbf{b} - \alpha \frac{dJ(\mathbf{b}, \mathbf{w}_1)}{d\mathbf{b}}$$

$$\text{temp1} = \mathbf{w}_1 - \alpha \frac{dJ(\mathbf{b}, \mathbf{w}_1)}{d\mathbf{w}_1}$$

$$\mathbf{b} = \text{temp0}$$

$$\mathbf{w}_1 = \text{temp1}$$




Incorrect form

$$\mathbf{b} = \mathbf{b} - \alpha \frac{dJ(\mathbf{b}, \mathbf{w}_1)}{d\mathbf{b}}$$

$$\mathbf{w}_1 = \mathbf{w}_1 - \alpha \frac{dJ(\mathbf{b}, \mathbf{w}_1)}{d\mathbf{w}_1}$$



Gradient Descent




خطوط مماس نشان داده شده دارای شیب یا مشتق مثبت هستند.
در نتیجه:

$$\frac{dJ(w^1)}{dw^1} > 0, \alpha > 0 \Rightarrow \alpha \frac{dJ(w^1)}{dw^1} > 0$$

$$\Rightarrow w^2 = w^1 - \alpha \frac{dJ(w^1)}{dw^1}$$

w کوچکتر میشود و به سمت چپ حرکت میکنیم.

Gradient Descent




خطوط مماس نشان داده شده دارای شیب یا مشتق مثبت هستند.
در نتیجه:

$$\frac{dJ(w^1)}{dw^1} < 0, \alpha > 0 \Rightarrow \alpha \frac{dJ(w^1)}{dw^1} < 0$$


$$\Rightarrow w^2 = w^1 - \alpha \frac{dJ(w^1)}{dw^1}$$

w بزرگتر میشود و به سمت راست حرکت میکنیم.

Choosing Learning Rate




α is too large




α is small

Gradient Descent Weakness



Gradient Descent



- 1) $\alpha > 0$

Repeat{

$$w = w - \alpha \frac{dJ(w)}{dw}$$

}until convergence

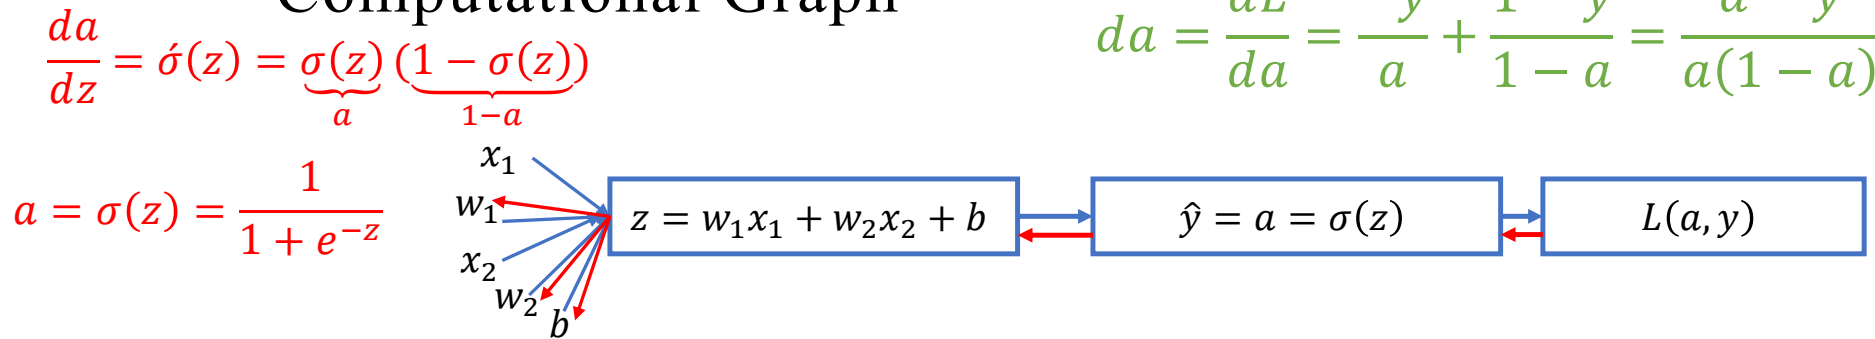
$$w = w - \alpha dw$$

$$z = w^T x + b$$

- $\hat{y} = a = \sigma(z)$
- $L(a, y) = -(y \log a + (1 - y) \log(1 - a))$

Gradient Descent

Computational Graph



$$L(a, y) = -(y \log a + (1 - y) \log(1 - a))$$

$$da = \frac{dL}{da} = \frac{-y}{a} + \frac{1 - y}{1 - a} = \frac{a - y}{a(1 - a)}$$

$$dz = \frac{dL}{dz} = \frac{dL(a, y)}{dz} = \frac{dL}{da} \times \frac{da}{dz} = \frac{a - y}{a(1 - a)} \times a(1 - a) = a - y$$

$$dw_1 = \frac{dL}{dw_1} = \frac{dL}{\underbrace{dz}_{\frac{dz}{dz}}} \times \underbrace{\frac{dz}{dw_1}}_{x_1} = x_1 dz$$

$$dw_2 = x_2 dz$$

$$db = \frac{dL}{db} = \frac{dL}{dz} \times \frac{\overset{1}{\widehat{dz}}}{db} = dz$$

Gradient Descent

- $J(w,b) = \frac{1}{m} \sum_{i=1}^m L(a^{(i)}, y^{(i)})$
- $a^{(i)} = \hat{y}^{(i)} = \sigma(z^{(i)}) = \sigma(wx^{(i)} + b)$

- $$dw_j = \frac{1}{m} \sum_{i=1}^m \frac{dL(a^{(i)}, y^{(i)})}{dw_j}$$

Logistic regression on m examples

$J = 0;$ $dw_1 = 0;$ $dw_2 = 0;$ $db = 0;$

$w_1 \leftarrow \text{random}$ $w_2 \leftarrow \text{random}$ $b \leftarrow \text{random}$

Repeat{

For $i=1$ **to** m

$$z^{(i)} = w^T x^{(i)} + b$$

$$a^{(i)} = \sigma(z^{(i)})$$

$$J += [y^{(i)} \text{Log} a^{(i)} + (1 - y^{(i)}) \text{Log}(1 - a^{(i)})]$$

$$dz^{(i)} = a^{(i)} - y^{(i)}$$

$$dw_1 += x_1^{(i)} dz^{(i)}$$

$$dw_2 += x_2^{(i)} dz^{(i)}$$

$$db += dz^{(i)}$$

$$J /= m;$$

$$dw_1 /= m;$$

$$dw_2 /= m;$$

$$db /= m;$$

$$w_1 = w_1 - \alpha dw_1 \quad w_2 = w_2 - \alpha dw_2$$

$$b = b - \alpha db$$

} until convergence

$$w^t = \begin{bmatrix} w_1^t \\ w_2^t \\ b^t \end{bmatrix} \quad w^{t+1} = \begin{bmatrix} w_1^{t+1} \\ w_2^{t+1} \\ b^{t+1} \end{bmatrix}$$

$$\|w^{t+1} - w^t\|_2 \leq \varepsilon$$

$$dw = \begin{bmatrix} dw_1 \\ dw_2 \\ db \end{bmatrix}$$

$$\|dw\| \leq \varepsilon = 10^{-4}$$

What's wrong with the code?

Logistic regression on m examples

$w_1 \leftarrow \text{random}$ $w_2 \leftarrow \text{random}$ $b \leftarrow \text{random}$

Repeat{

$J = 0;$ $dw_1 = 0;$ $dw_2 = 0;$ $db = 0;$

For $i=1$ *to* m

$$z^{(i)} = w^T x^{(i)} + b$$

$$a^{(i)} = \sigma(z^{(i)})$$

$$J += [y^{(i)} \text{Log} a^{(i)} + (1 - y^{(i)}) \text{Log}(1 - a^{(i)})]$$

$$dz^{(i)} = a^{(i)} - y^{(i)}$$

$$dw_1 += x_1^{(i)} dz^{(i)}$$

$$dw_2 += x_2^{(i)} dz^{(i)}$$

$$db += dz^{(i)}$$

$J /= m;$

$dw_1 /= m;$

$dw_2 /= m;$

$db /= m;$

$$w_1 = w_1 - \alpha dw_1 \quad w_2 = w_2 - \alpha dw_2$$

$$b = b - \alpha db$$

} until convergence

$$w^t = \begin{bmatrix} w_1^t \\ w_2^t \\ b^t \end{bmatrix} \quad w^{t+1} = \begin{bmatrix} w_1^{t+1} \\ w_2^{t+1} \\ b^{t+1} \end{bmatrix}$$

$$\|w^{t+1} - w^t\|_2 \leq \varepsilon$$

$$dw = \begin{bmatrix} dw_1 \\ dw_2 \\ db \end{bmatrix}$$

$$\|dw\| \leq \varepsilon = 10^{-4}$$

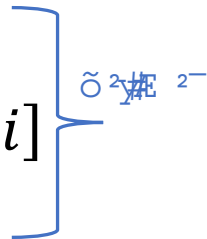
Basics of Neural Network Programming

Vectorizing Logistic Regression's Gradient Computation

Vectorizing Logistic Regression

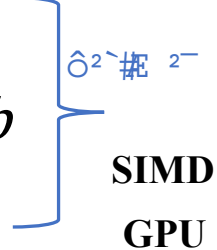
- $z=0;$

$\text{For } i \text{ in range}(n_x)$
 $z += w[i] * x[i]$
 $z += b$



- $z=0;$

$z = \text{np.dot}(w, x) + b$



SIMD
GPU

Vectorizing Logistic Regression

$$\begin{aligned} z^{(1)} &= w^T x^{(1)} + b & z^{(2)} &= w^T x^{(2)} + b \\ a^{(1)} &= \sigma(z^{(1)}) & a^{(2)} &= \sigma(z^{(2)}) \end{aligned}$$

$$\begin{aligned} z^{(m)} &= w^T x^{(m)} + b \\ a^{(m)} &= \sigma(z^{(m)}) \end{aligned}$$

$$\bullet X = \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & & | \end{bmatrix} \in R^{n_x \times m} \quad \underbrace{[w_1 \dots w_{n_x}]}_{W^T} \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & & | \end{bmatrix}$$

$$\bullet Z = \begin{bmatrix} \underbrace{z^{(1)}}_{w^T x^{(1)} + b} & \underbrace{z^{(2)}}_{w^T x^{(2)} + b} & \dots & \underbrace{z^{(m)}}_{w^T x^{(m)} + b} \end{bmatrix} = w^T X + [b \ b \ \dots \ b]_{1 \times m}$$

$$\bullet Z = \underbrace{np \cdot dot(w \cdot T, X)}_{1 \times m} + \underbrace{b}_{(1,1)} \text{ "broadcasting" } [b, b \dots, b]_{1 \times m}$$

Vectorizing Logistic Regression

- $A = [a^{(1)}, a^{(2)}, \dots, a^{(m)}] = \sigma(\underbrace{\underline{Z}}_{1 \times m})$

- $dz^{(1)} = a^{(1)} - y^{(1)} \quad dz^{(2)} = a^{(2)} - y^{(2)}$

- $dZ = [dz^{(1)} \ dz^{(2)} \ \dots \ dz^{(m)}]_{1 \times m}$

- $A = [a^{(1)} \ a^{(2)} \ \dots \ a^{(m)}] \quad Y = [y^{(1)} \ y^{(2)} \ \dots \ y^{(m)}]$

- $dZ = A - Y = [a^{(1)} - y^{(1)} \ a^{(2)} - y^{(2)} \ \dots \ a^{(m)} - y^{(m)}]$

- $\begin{bmatrix} dw_1 \\ dw_2 \\ dw_3 \\ \vdots \\ dw_{nx} \end{bmatrix}_{nx \times 1} \quad \begin{cases} dw=0 \\ dw += \overbrace{\tilde{x}^T}^{\text{عدد بردار}} \overbrace{\tilde{dz}^1}^{a^{(1)}-y^{(1)}} \\ dw += x^2 dz^2 \\ dw += x^m dz^m \\ dw/=m \end{cases} \quad \begin{cases} db=0 \\ db += \overbrace{\tilde{dz}^1}^{a^{(1)}-y^{(1)}} \\ db += dz^2 \\ db += dz^m \\ db/=m \end{cases} \quad db = \frac{1}{m} \sum_{i=1}^m dz^{(i)} = \frac{1}{m} np.sum(dz) \quad db = \frac{1}{m} dz \times 1$

$$1 = \begin{bmatrix} 1 \\ 1 \\ 1 \\ \vdots \\ \vdots \\ \vdots \\ 1 \end{bmatrix}_{m \times 1}$$

$$dw = \frac{1}{m} [x^{(1)} dz^1 + x^{(2)} dz^2 + \dots + x^{(m)} dz^m]$$

$$\boxed{dw = \frac{1}{m} X dz^T} = \frac{1}{m} \begin{bmatrix} x^{(1)} x^{(2)} & \dots & \dots & x^{(m)} \end{bmatrix} \times \begin{bmatrix} dz^1 \\ dz^2 \\ dz^3 \\ \vdots \\ \vdots \\ dz^m \end{bmatrix}$$

Vectorizing Logistic Regression

- $w_1, w_2, b \leftarrow \text{random}$

For iter in range(1000)

1 epoch

$$Z = W^T X + b = np \cdot \text{dot}(w.T, X) + b$$

$$A = \sigma(Z)$$

$$dZ = A - Y$$

$$dw = \frac{1}{m} X dz^T$$

$$db = \frac{1}{m} np.sum(dz)$$

$$w = w - \alpha dw$$

$$b = b - \alpha db$$

$$w = np \cdot \text{random} \cdot \text{randn}(n_x, 1)$$

Logistic Regression Cost function

- $\hat{y} = \sigma(w^T x + b) \quad 0 < \sigma(z) = \frac{1}{1+e^{-z}} < 1$
- $\hat{y} = p(y = 1|x)$
- $\left. \begin{array}{l} \text{if } y = 1 : p(y|x) = \hat{y} \\ \text{if } y = 0 : p(y|x) = 1 - \hat{y} \end{array} \right\} p(y|x)$
- $\boxed{p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}}$ *distribution? Bernoulli*
- $\text{if } y = 1 : p(y|x) = \hat{y}$
- $\text{if } y = 0 : p(y|x) = 1 - \hat{y}$
- $\log p(y|x) = \log[\hat{y}^y (1 - \hat{y})^{1-y}] = y \log \hat{y} + (1 - y) \log(1 - \hat{y})$
 $= -L(\hat{y}, y) \text{ Max Likelihood}$

Logistic Regression Cost function

- $\log P(\text{labels in trainingset}) = \log \prod_{i=1}^m P(y^i | x^i)$
- $\log P(\dots) = \sum_{i=1}^m \log P(y^{(i)} | x^{(i)}) = -\sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$

- Cost function
$$\underbrace{J(w, b)}_{\text{minimize}} = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$$

Neural Networks

