



Deep Learning

Logistic Regression Classifier

Dr. Mehran Safayani



safayani@iut.ac.ir



safayani.iut.ac.ir



<https://www.aparat.com/mehran.safayani>



https://github.com/safayani/deep_learning_course



Supervised Learning

Input(x)	Output (y)	Application
Home features	Price	Real Estate
Ad, user info	Click on ad? (0/1)	Online Advertising
Image	Object (1,...,1000)	Photo tagging
Audio	Text transcript	Speech recognition
English	Chinese	Machine translation
Image, Radar info	Position of other cars	Autonomous driving

Basics of Neural Network Programming

Binary Classification

Binary Classification



—————→ 1 (cat) vs 0 (non cat)

		Blue			
Green		255	134	93	22
Red	255	134	202	22	2
	255	231	42	22	4
	123	94	83	2	192
	34	44	187	92	34
	34	76	232	124	94
	67	83	194	202	

Binary classification

[illegible]

$$\vec{x} = \begin{bmatrix} 255 \\ 231 \\ \vdots \\ 254 \\ 253 \\ 250 \\ 220 \end{bmatrix} \quad 64 \times 64 \times 3 = \underbrace{12288}_{n_x}$$

$$\vec{x} \rightarrow \text{model} \rightarrow \hat{y}$$

- Notation

$$(\vec{x}, y) \quad x \in R^{n_x}, y \in \{0,1\}$$

Binary classification

- m training example: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

$$\bullet X = \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & & | \end{bmatrix} \begin{matrix} \updownarrow \\ n_x \end{matrix}$$

$\xleftarrow{\quad m \quad} \xrightarrow{\quad}$

$$X \in \mathbb{R}^{n_x \times m}$$

$$X.\text{shape} = (n_x, m)$$

$$Y = [y^{(1)}, y^{(2)}, \dots, y^{(m)}]$$

$$Y \in \mathbb{R}^{1 \times m}$$

$$Y.\text{shape} = (1, m)$$

Logistic Regression

- Given \mathbf{x} , output $\hat{y} = P(y=1|\mathbf{x})$ $0 \leq \hat{y} \leq 1$


$\mathbf{x} \in \mathbb{R}^{n_x}$ parameters: $\mathbf{w} \in \mathbb{R}^{n_x}$, $b \in \mathbb{R}$

$$\hat{y} = \mathbf{w}^T \mathbf{x} + b$$

Sigmoid function $\sigma(z) = \frac{1}{1+e^{-z}}$

if z large $\sigma(z) \approx 1$

if z large negative $\sigma(z) \approx 0$



Logistic Regression

$$\bullet \hat{y} = \sigma(\underbrace{w^T x + b}_z) \quad \hat{y} = w^T x$$

$$x_0 = 1, x \in \mathbb{R}^{n_x+1}$$

$$W = \begin{bmatrix} b = w_0 \\ w_1 \\ w_2 \\ w_3 \\ \vdots \\ \vdots \\ \vdots \\ w_{n_x} \end{bmatrix} \begin{matrix} \left. \vphantom{\begin{bmatrix} b = w_0 \\ w_1 \\ w_2 \\ w_3 \\ \vdots \\ \vdots \\ \vdots \\ w_{n_x} \end{bmatrix}} \right\} b \\ \left. \vphantom{\begin{bmatrix} w_1 \\ w_2 \\ w_3 \\ \vdots \\ \vdots \\ \vdots \\ w_{n_x} \end{bmatrix}} \right\} w \end{matrix}$$

Logistic Regression cost function

- Loss (error) function: $L(\hat{y}, y) = \frac{1}{2} (\hat{y} - y)^2 = \frac{1}{2} (\sigma(w^T x + b) - y)^2$ SE: Square Error

- What is the problem?



- .

Convexity

Function $h(u)$ with $u \in X$ is **convex** if for any $u, v \in X$ and for any $0 \leq \lambda \leq 1$ we have:

$$h(\lambda u + (1 - \lambda)v) \leq \lambda h(u) + (1 - \lambda) h(v)$$

برای توابع محدب هر بهینه محلی یک بهینه سراسری است.




Cross Entropy

- $L(\hat{y}, y) = -(y \log \hat{y} + (1-y) \log(1-\hat{y}))$

if $y=1$: $L(\hat{y}, y) = -\log \hat{y}$

if $y=0$: $L(\hat{y}, y) = -\log(1-\hat{y})$




<https://datamonje.com/classification-loss-functions/>

- Cost function:
$$J(w,b) = \frac{1}{m} \sum_{i=1}^m L(y^{(i)}, \hat{y}^{(i)})$$
$$= -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1-y^{(i)}) \log(1-\hat{y}^{(i)})]$$

Decision Boundary

$$\theta = \begin{bmatrix} -3 \\ 1 \\ 1 \end{bmatrix}$$


$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$



Predict $y = 1, \text{ if } \underbrace{-3 + x_1 + x_2}_{\theta^T x} \geq 0$

$x_1 + x_2 \geq 3$

Non-Linear Decision Boundaries



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_1^2 x_2 + \theta_5 x_1^2 x_2^2 + \theta_6 x_1^3 x_2 + \dots)$$


$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

\downarrow \downarrow \downarrow
 -1 0 0

Predict $y = 1, \text{ if } \underbrace{-3 + x_1^2 + x_2^2}_{\theta^T x} \geq 0$

$$x_1^2 + x_2^2 \geq 1$$

$$\theta = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}$$



Gradient Descent

Cost Function

Minimize $J(\mathbf{b}, w_1)$

\mathbf{b}, w_1

If $J(w_1) = (w_1 - 2)^2$

$$\frac{dJ(w_1)}{dw_1} = 0$$



$$\frac{dJ(w_1)}{dw_1} = 2(w_1 - 2) = 0$$



$$w_1 = 2$$

Gradient Descent

Minimize $J(\mathbf{b}, w_1)$

\mathbf{b}, w_1

Minimize $J(\mathbf{b}, w_1, \dots, w_n)$

$\mathbf{b}, w_1, \dots, w_n$

Repeat until convergence: {

For $j=0, \dots, n$

$b = w_0$

$$w_j = w_j - \alpha \frac{dJ(\mathbf{b}, w_1, \dots, w_n)}{dw_j}$$

}

α is **learning rate**

Updating all w_j *Simultaneously*

Convergence condition:

$$\|W^{t+1} - W^t\|_2 \leq \varepsilon$$

Gradient Descent

Correct form

$$\text{temp0} = \mathbf{b} - \alpha \frac{dJ(\mathbf{b}, \mathbf{w}_1)}{d\mathbf{b}}$$

$$\text{temp1} = \mathbf{w}_1 - \alpha \frac{dJ(\mathbf{b}, \mathbf{w}_1)}{d\mathbf{w}_1}$$

$$\mathbf{b} = \text{temp0}$$

$$\mathbf{w}_1 = \text{temp1}$$



بروزرسانی همزمان


Incorrect form

$$\mathbf{b} = \mathbf{b} - \alpha \frac{dJ(\mathbf{b}, \mathbf{w}_1)}{d\mathbf{b}}$$

$$\mathbf{w}_1 = \mathbf{w}_1 - \alpha \frac{dJ(\mathbf{b}, \mathbf{w}_1)}{d\mathbf{w}_1}$$



Gradient Descent




خطوط مماس نشان داده شده دارای شیب یا مشتق مثبت هستند.
در نتیجه:

$$\frac{dJ(w^1)}{dw^1} > 0, \alpha > 0 \Rightarrow \alpha \frac{dJ(w^1)}{dw^1} > 0$$

$$\Rightarrow w^2 = w^1 - \alpha \frac{dJ(w^1)}{dw^1}$$

w کوچکتر میشود و به سمت چپ حرکت میکنیم.

Gradient Descent




خطوط مماس نشان داده شده دارای شیب یا مشتق مثبت هستند.
در نتیجه:

$$\frac{dJ(w^1)}{dw^1} < 0, \alpha > 0 \Rightarrow \alpha \frac{dJ(w^1)}{dw^1} < 0$$


$$\Rightarrow w^2 = w^1 - \alpha \frac{dJ(w^1)}{dw^1}$$

w بزرگتر میشود و به سمت راست حرکت میکنیم.

Choosing Learning Rate




α is too large

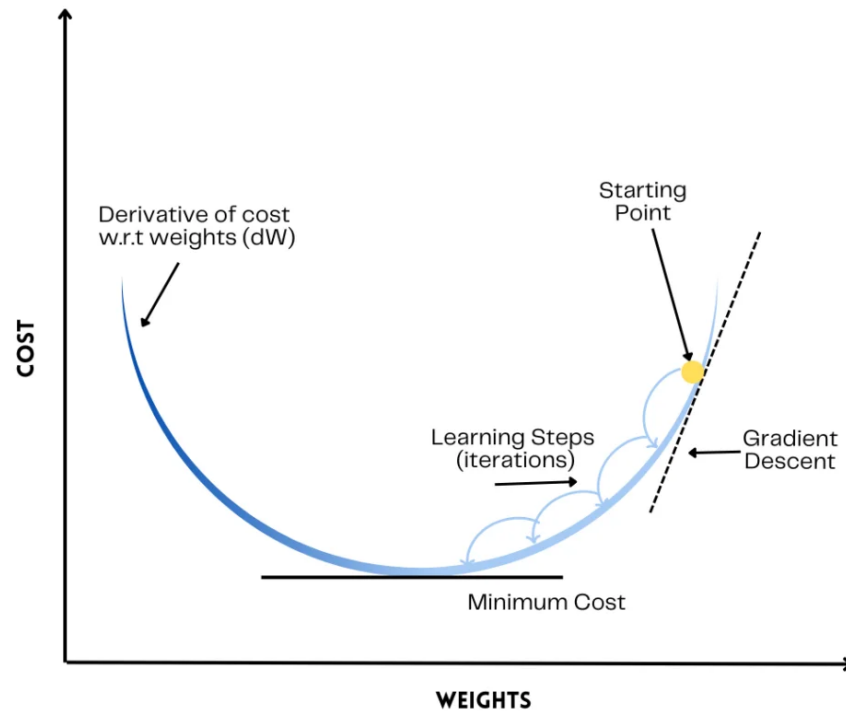


α is small

Gradient Descent Weakness



Gradient Descent



- 1) $\alpha > 0$

Repeat{

$$w = w - \alpha \underbrace{\frac{dJ(w)}{dw}}$$

}until convergence

$$w = w - \alpha dw$$

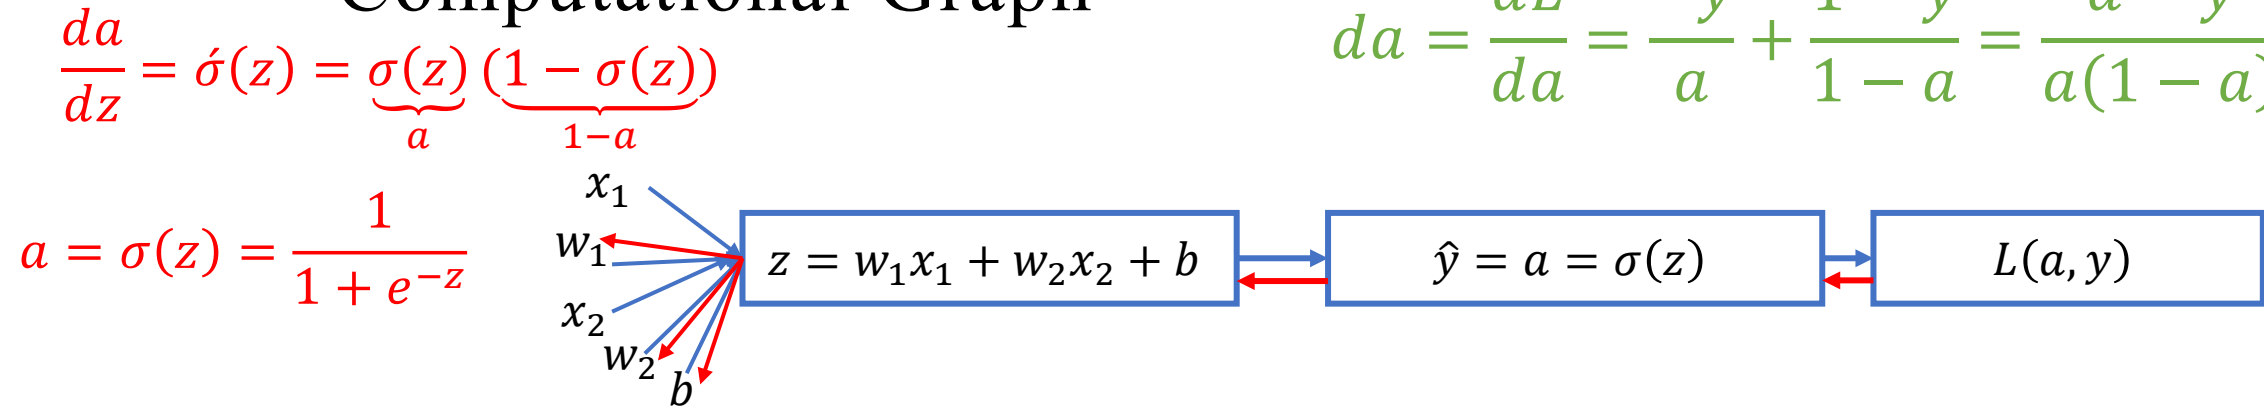
$$z = w^T x + b$$

- $\hat{y} = a = \sigma(z)$
- $L(a, y) = -(y \log a + (1 - y) \log(1 - a))$

Gradient Descent

$$L(a, y) = -(y \log a + (1 - y) \log(1 - a))$$

Computational Graph



$$dz = \frac{dL}{dz} = \frac{dL(a, y)}{dz} = \frac{dL}{da} \times \frac{da}{dz} = \frac{a - y}{a(1 - a)} \times a(1 - a) = a - y$$

$$dw_1 = \frac{dL}{dw_1} = \underbrace{\frac{dL}{dz}}_{dz} \times \underbrace{\frac{dz}{dw_1}}_{x_1} = x_1 dz \quad dw_2 = x_2 dz \quad db = \frac{dL}{db} = \frac{dL}{dz} \times \underbrace{\frac{dz}{db}}_1 = dz$$

Gradient Descent

- $J(w,b) = \frac{1}{m} \sum_{i=1}^m L(a^{(i)}, y^{(i)})$

- $a^{(i)} = \hat{y}^{(i)} = \sigma(z^{(i)}) = \sigma(wx^{(i)} + b)$

- $dw_j = \frac{1}{m} \sum_{i=1}^m \frac{dL(a^{(i)}, y^{(i)})}{dw_j}$

Logistic regression on m examples

$J = 0;$ $dw_1 = 0;$ $dw_2 = 0;$ $db = 0;$

$w_1 \leftarrow \text{random}$ $w_2 \leftarrow \text{random}$ $b \leftarrow \text{random}$

Repeat{

For $i=1$ **to** m

$$z^{(i)} = w^T x^{(i)} + b$$

$$a^{(i)} = \sigma(z^{(i)})$$

$$J += [y^{(i)} \text{Log} a^{(i)} + (1 - y^{(i)}) \text{Log}(1 - a^{(i)})]$$

$$dz^{(i)} = a^{(i)} - y^{(i)}$$

$$dw_1 += x_1^{(i)} dz^{(i)}$$

$$dw_2 += x_2^{(i)} dz^{(i)}$$

$$db += dz^{(i)}$$

$$J /= m;$$

$$dw_1 /= m;$$

$$dw_2 /= m;$$

$$db /= m;$$

$$w_1 = w_1 - \alpha dw_1$$

$$w_2 = w_2 - \alpha dw_2$$

$$b = b - \alpha db$$

} until convergence

$$w^t = \begin{bmatrix} w_1^t \\ w_2^t \\ b^t \end{bmatrix} \quad w^{t+1} = \begin{bmatrix} w_1^{t+1} \\ w_2^{t+1} \\ b^{t+1} \end{bmatrix}$$

$$\|w^{t+1} - w^t\|_2 \leq \varepsilon$$

$$dw = \begin{bmatrix} dw_1 \\ dw_2 \\ db \end{bmatrix}$$

$$\|dw\| \leq \varepsilon = 10^{-4}$$

What's wrong with the code?

Logistic regression on m examples

$w_1 \leftarrow \text{random}$ $w_2 \leftarrow \text{random}$ $b \leftarrow \text{random}$

Repeat{

$J = 0;$ $dw_1 = 0;$ $dw_2 = 0;$ $db = 0;$

For $i=1$ to m

$$z^{(i)} = w^T x^{(i)} + b$$

$$a^{(i)} = \sigma(z^{(i)})$$

$$J += [y^{(i)} \text{Log} a^{(i)} + (1 - y^{(i)}) \text{Log}(1 - a^{(i)})]$$

$$dz^{(i)} = a^{(i)} - y^{(i)}$$

$$dw_1 += x_1^{(i)} dz^{(i)}$$

$$dw_2 += x_2^{(i)} dz^{(i)}$$

$$db += dz^{(i)}$$

$$J /= m;$$

$$dw_1 /= m;$$

$$dw_2 /= m;$$

$$db /= m;$$

$$w_1 = w_1 - \alpha dw_1$$

$$w_2 = w_2 - \alpha dw_2$$

$$b = b - \alpha db$$

} until convergence

$$w^t = \begin{bmatrix} w_1^t \\ w_2^t \\ b^t \end{bmatrix} \quad w^{t+1} = \begin{bmatrix} w_1^{t+1} \\ w_2^{t+1} \\ b^{t+1} \end{bmatrix}$$

$$\|w^{t+1} - w^t\|_2 \leq \varepsilon$$

$$dw = \begin{bmatrix} dw_1 \\ dw_2 \\ db \end{bmatrix}$$

$$\|dw\| \leq \varepsilon = 10^{-4}$$

Vectorizing Logistic Regression

- $z=0;$

For i *in* $\text{range}(n_x)$
 $z += w[i] * x[i]$
 $z += b$

} First method

- $z=0;$

$z = \text{np.dot}(w, x) + b$

} Second method
SIMD
GPU

Vectorizing Logistic Regression

$$z^{(1)} = w^T x^{(1)} + b \quad z^{(2)} = w^T x^{(2)} + b$$

$$a^{(1)} = \sigma(z^{(1)}) \quad a^{(2)} = \sigma(z^{(2)})$$

$$z^{(m)} = w^T x^{(m)} + b$$

$$a^{(m)} = \sigma(z^{(m)})$$

$$\bullet X = \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & & | \end{bmatrix} \in R^{n_x \times m} \quad \underbrace{[w_1 \dots w_{n_x}]}_{W^T} \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & & | \end{bmatrix}$$

$$\bullet Z = \begin{bmatrix} \underbrace{z^{(1)}}_{w^T x^{(1)} + b} & \underbrace{z^{(2)}}_{w^T x^{(2)} + b} & \dots & \underbrace{z^{(m)}}_{w^T x^{(m)} + b} \end{bmatrix} = w^T X + [b \ b \ \dots \ b]_{1 \times m}$$

$$\bullet Z = \underbrace{np \cdot dot(w \cdot T, X)}_{1 \times m} + \underbrace{b}_{(1,1)} \text{ "broadcasting" } [b, b \dots, b]_{1 \times m}$$

Vectorizing Logistic Regression

- $A = [a^{(1)}, a^{(2)}, \dots, a^{(m)}] = \sigma(\underbrace{Z}_{1 \times m})$

- $dz^{(1)} = a^{(1)} - y^{(1)} \quad dz^{(2)} = a^{(2)} - y^{(2)}$

- $dZ = [dz^{(1)} \ dz^{(2)} \ \dots \ dz^{(m)}]_{1 \times m}$

- $A = [a^{(1)} \ a^{(2)} \ \dots \ a^{(m)}] \quad Y = [y^{(1)} \ y^{(2)} \ \dots \ y^{(m)}]$

- $dZ = A - Y = [a^{(1)} - y^{(1)} \ a^{(2)} - y^{(2)} \ \dots \ a^{(m)} - y^{(m)}]$

- $\begin{bmatrix} dw_1 \\ dw_2 \\ dw_3 \\ \vdots \\ dw_{nx} \end{bmatrix}_{nx \times 1}$

- $\begin{cases} dw=0 \\ dw += \overset{\text{عدد}}{\tilde{x}^1} \overset{\text{عدد}}{\tilde{dz}^1} \\ dw += x^2 dz^2 \\ dw += x^m dz^m \\ dw /= m \end{cases}$

- $\begin{cases} db=0 \\ db += \overset{\text{عدد}}{\tilde{dz}^1} \\ db += dz^2 \\ db += dz^m \\ db /= m \end{cases}$

$$db = \frac{1}{m} \sum_{i=1}^m dz^{(i)} = \frac{1}{m} np.sum(dz)$$

$$db = \frac{1}{m} dz \times 1$$

$$1 = \begin{bmatrix} 1 \\ 1 \\ 1 \\ \vdots \\ \vdots \\ \vdots \\ 1 \end{bmatrix}_{m \times 1}$$

$$dz = [dz^1 \ dz^2 \ \dots \ dz^m]_{1 \times m}$$

$$dw = \frac{1}{m} [x^{(1)} dz^1 + x^{(2)} dz^2 + \dots + x^{(m)} dz^m]$$

$$\boxed{dw = \frac{1}{m} X dz^T} = \frac{1}{m} \begin{bmatrix} x^{(1)} x^{(2)} & \dots & \dots & x^{(m)} \end{bmatrix} \times \begin{bmatrix} dz^1 \\ dz^2 \\ dz^3 \\ \vdots \\ \vdots \\ dz^m \end{bmatrix}$$

Vectorizing Logistic Regression

- $w_1, w_2, b \leftarrow \text{random}$

For iter in range(1000)

1 epoch

$$Z = W^T X + b = np \cdot \text{dot}(w.T, X) + b$$

$$A = \sigma(Z)$$

$$dZ = A - Y$$

$$dw = \frac{1}{m} X dz^T$$

$$db = \frac{1}{m} np.sum(dz)$$

$$w = w - \alpha dw$$

$$b = b - \alpha db$$

$$w = np \cdot \text{random} \cdot \text{randn}(n_x, 1)$$

Logistic Regression Cost function

- $\hat{y} = \sigma(w^T x + b) \quad 0 < \sigma(z) = \frac{1}{1+e^{-z}} < 1$
- $\hat{y} = p(y = 1|x)$
- $\left. \begin{array}{l} \text{if } y = 1 : p(y|x) = \hat{y} \\ \text{if } y = 0 : p(y|x) = 1 - \hat{y} \end{array} \right\} p(y|x)$
- $\boxed{p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}}$ *distribution? Bernoulli*
- $\text{if } y = 1 : p(y|x) = \hat{y}$
- $\text{if } y = 0 : p(y|x) = 1 - \hat{y}$
- $\log p(y|x) = \log[\hat{y}^y (1 - \hat{y})^{1-y}] = y \log \hat{y} + (1 - y) \log(1 - \hat{y})$
 $= -L(\hat{y}, y) \text{ Max Likelihood}$

Logistic Regression Cost function

- $\log P(\text{labels in trainingset}) = \log \prod_{i=1}^m P(y^i | x^i)$
- $\log P(\cdots) = \sum_{i=1}^m \log P(y^{(i)} | x^{(i)}) = -\sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$
- Cost function
$$\underbrace{J(w, b)}_{\text{minimize}} = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$$

Neural Networks

